Test Case / Original Request

Run PTC-005 using Jmeter. Execute the real performance test and collect monitoring results. Assess the security surface for the same job and exact configured target.Analyze the current job's results and generate the final PDF with security coverage and evidence, plus links to the native JMeter report, raw results, and final PDF.

REAL Execution Outcome

Execution status: COMPLETED
Runner status: completed
Outcome details: REAL JMeter execution completed. Configured thresholds passed.
Process exit code: 0
Thresholds passed: True

Executive Summary

Application: Unknown
Scenario: Task Status API — Load Check
Job ID: JOB-5d2d8f4c8b924f0f8d456ab6eea88c92
Overall Result: WARNING

Performance test execution completed and passed primary thresholds, but outlier max response times were seen for /todos/1, similar to prior runs of this scenario. System memory on the load generator exceeded 92%, flagged as a CRITICAL monitoring issue, though client-observed error rate was zero and average response times were satisfactory. No security analysis was performed for this run. Periodic long-tail latency spikes continue to merit focused investigation, with supporting historical evidence and a relevant remediation runbook available.

Quick Links

Artifact	*Open*
Native HTML Dashboard	Open →
Dashboard + Raw Results (ZIP)	Open →
Raw Results	Open →

Execution Profile and Artifacts

Execution mode: REAL
Workload profile:
- tool: jmeter
- target_url: https://jsonplaceholder.typicode.com
- endpoint: /todos/1
- method: GET
- users: 25

- max_rps: 15
- duration_seconds: 60
- ramp_up_seconds: 5
- expected_status_code: 200
- p95_threshold_ms: 5000
Artifact directory: src/generated_scripts/Jmeter native reports/JOB-5d2d8f4c8b924f0f8d456ab6eea88c92
Artifact references:
- aligned_time_series: src/generated_scripts/Jmeter native reports/JOB-5d2d8f4c8b924f0f8d456ab6eea88c92/jmeter_PTC_005_plan_aligned_time_series.csv
- generated_script: src/generated_scripts/Jmeter native reports/JOB-5d2d8f4c8b924f0f8d456ab6eea88c92/jmeter_PTC_005_plan.jmx
- jmeter_log: src/generated_scripts/Jmeter native reports/JOB-5d2d8f4c8b924f0f8d456ab6eea88c92/jmeter_PTC_005_plan_jmeter.log
- measurement_metadata: src/generated_scripts/Jmeter native reports/JOB-5d2d8f4c8b924f0f8d456ab6eea88c92/jmeter_PTC_005_plan_measurement_metadata.json
- monitoring_samples: src/generated_scripts/Jmeter native reports/JOB-5d2d8f4c8b924f0f8d456ab6eea88c92/jmeter_PTC_005_plan_runtime_samples.csv
- Native JMeter HTML dashboard
- JMeter dashboard + raw JTL ZIP
- native_report_entrypoint: src/generated_scripts/Jmeter native reports/JOB-5d2d8f4c8b924f0f8d456ab6eea88c92/jmeter_PTC_005_plan_native_report.html
- normalized_metrics: src/generated_scripts/Jmeter native reports/JOB-5d2d8f4c8b924f0f8d456ab6eea88c92/jmeter_PTC_005_plan_normalized_metrics.json
- Raw JMeter results (JTL)
- runner_log: src/generated_scripts/Jmeter native reports/JOB-5d2d8f4c8b924f0f8d456ab6eea88c92/jmeter_PTC_005_plan_runner.log
- time_series: src/generated_scripts/Jmeter native reports/JOB-5d2d8f4c8b924f0f8d456ab6eea88c92/jmeter_PTC_005_plan_time_series.csv
The interactive native HTML dashboard is linked as a separate artifact and is not embedded in this PDF.

Time-Series Charts (REAL Measurements)

Charts show measured values in five-second aligned buckets. Host CPU is from the load-generator machine, not the target.

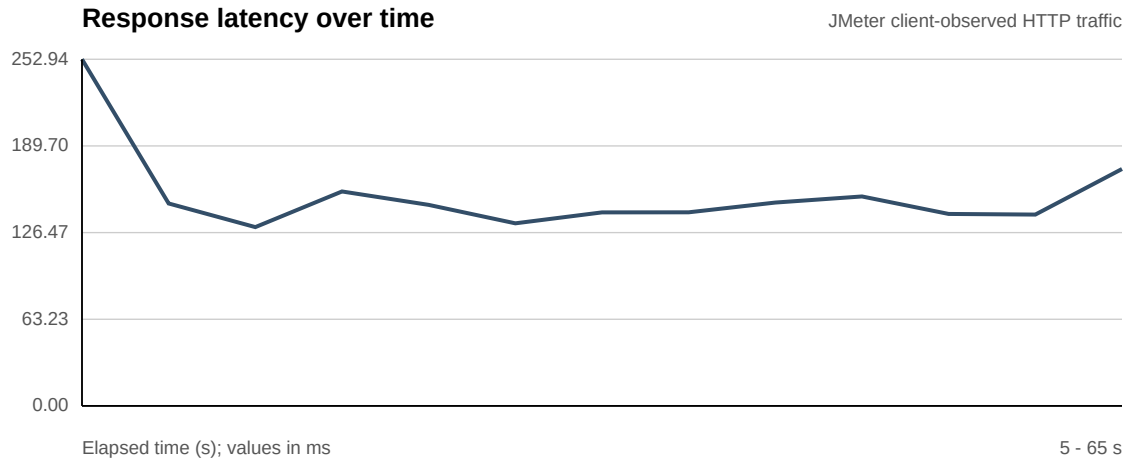

Throughput over time

JMeter client-observed HTTP traffic

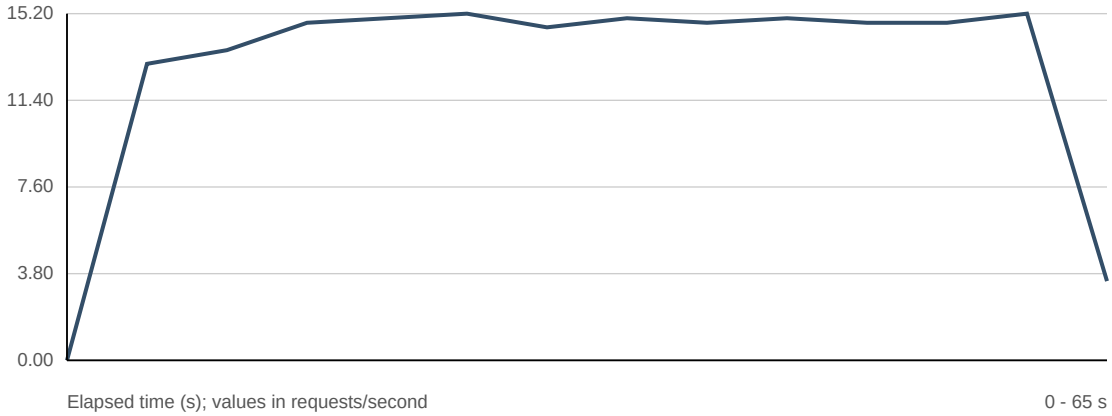

Load-generator host CPU over time

psutil load-generator host

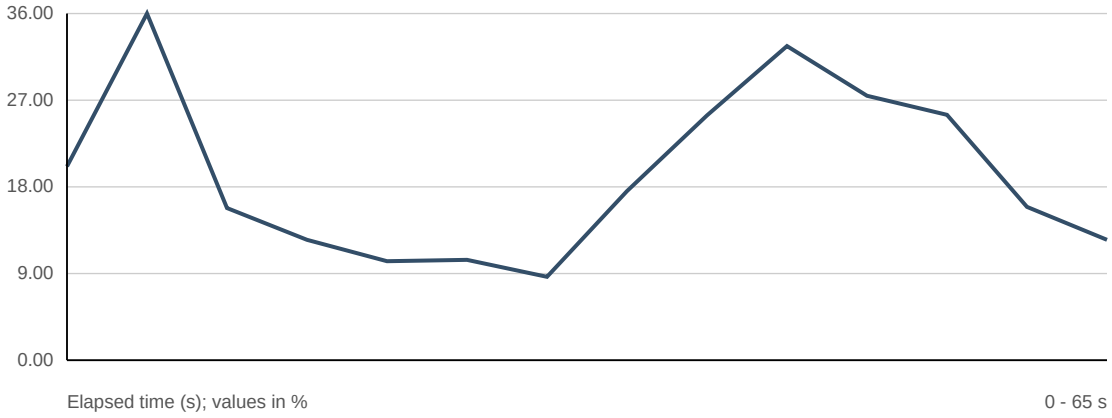

Agent Activity Summary

Agent	*Summary*
Test Execution	Load test executed on /todos/1 (Unknown - Task Status API — Load Check). Results: 887 total requests, 887 successful, 0 failed. Avg response time: 151.74295377677564 ms; 95th percentile: 267.0 ms; Max: 999.0 ms. Throughput: 13.179244518638933 TPS; Error rate: 0.0%. Latency percentiles (ms): p50=143.0, p90=206.0, p95=267.0, p99=489.0

Agent	*Summary*
Monitoring	CPU usage: 19.64% average, 58.8% peak; memory: 92.87% average, 95.5% peak; DB connections: not collected, TPS: 13.179244518638933, Response time: 151.74295377677564 ms. Data source and scope: REAL: psutil load-generator host samples and jmeter client-observed metrics. Measurement provenance: cpu_usage=psutil (load_generator_host); cpu_peak_usage=psutil (load_generator_host); memory_usage=psutil (load_generator_host); memory_peak_usage=psutil (load_generator_host); response_time=jmeter (load_generator_observed_http_traffic); tps=jmeter (load_generator_observed_http_traffic); db_connections=not_collected (No target-specific database telemetry provider is configured.); total_requests=jmeter JTL (load_generator_observed_http_traffic); failed_requests=jmeter JTL (load_generator_observed_http_traffic); error_rate_pct=jmeter JTL (load_generator_observed_http_traffic); latency_percentiles_ms=jmeter JTL (load_generator_observed_http_traffic). Alerts triggered: ['High load-generator host memory usage']
Security Validation	Status: NOT ASSESSED. Security score: N/A. Live security check: not assessed. Issue: A matching job-scoped security result was unavailable. Assessment limitations: - No matching job-scoped security result was available to this report.
Analysis	Performance test execution completed and passed primary thresholds, but outlier max response times were seen for /todos/1, similar to prior runs of this scenario. System memory on the load generator exceeded 92%, flagged as a CRITICAL monitoring issue, though client-observed error rate was zero and average response times were satisfactory. No security analysis was performed for this run. Periodic long-tail latency spikes continue to merit focused investigation, with supporting historical evidence and a relevant remediation runbook available.
Reporting	A PDF report was generated, including all findings, recommendations, and an executive summary.

Per-API Findings (Test Execution)

- /todos/1 [WARNING / medium]
avg=151.74295377677564ms p95=267.0ms max=999.0ms throughput=13.179244518638933 req/s error_rate=0.0% http_status=200
* Max response time (999.0ms) is 3.7x p95 (267.0ms) -- outlier spikes present.
history: Previous runs of API /todos/1 under the same scenario have also shown max response time outliers (e.g., prior max 1413ms with p95 of 385ms, giving a similar 3.7x ratio). This indicates long-tail latency has been a recurring characteristic for this test profile, not a regression versus earlier measured performance.
recommendation: Investigate the cause of sporadic high-latency spikes for /todos/1, as this pattern is consistent with prior runs and might reflect backend or host-level contention. Consider referencing the existing runbook for high-latency APIs. Due to the similar 3.7x p95-max pattern (as previously observed), focus should be on isolating resources or code paths hit by outlier requests. No recent SLA or postmortem incident documents were directly matched to this endpoint.

Infrastructure Health (Monitoring)

Status: WARNING (monitoring_agent reported: CRITICAL)
CPU usage: 19.64% Memory usage: 92.87% DB connections: None TPS: 13.179244518638933
Response time (monitoring): 151.74295377677564ms
- Memory usage 92.87% exceeds 85.0% threshold.
- 1 monitoring alert(s) reported: ['High load-generator host memory usage']
Monitoring alerts: ['High load-generator host memory usage']

Security Findings (Security Agent)

Status: NOT ASSESSED Security score: N/A
Security result job ID: JOB-5d2d8f4c8b924f0f8d456ab6eea88c92.
Live security check: not assessed.
Issues:
- A matching job-scoped security result was unavailable.
Assessment limitations:
- No matching job-scoped security result was available to this report.

Recommended Next Steps

- /todos/1: Investigate the cause of sporadic high-latency spikes for /todos/1, as this pattern is consistent with prior runs and might reflect backend or host-level contention. Consider referencing the existing runbook for high-latency APIs. Due to the similar 3.7x p95-max pattern (as previously observed), focus should be on isolating resources or code paths hit by outlier requests. No recent SLA or postmortem incident documents were directly matched to this endpoint.
- Infrastructure: review CPU/memory/DB connection anomalies above.