Test Case / Original Request

Run test case PTC-003 using Jmeter, monitor it, analyze the results, and generate the final PDF report with links to the native JMeter report and raw results.

REAL Execution Outcome

Execution status: COMPLETED
Runner status: completed
Outcome details: REAL JMeter execution completed. Configured thresholds passed.
Process exit code: 0
Thresholds passed: True

Executive Summary

Application: Unknown
Scenario: Content Feedback API — Load Check
Job ID: JOB-43d9c3eab8984222819f9c2889df9702
Overall Result: WARNING

Test execution completed successfully with no failed requests and healthy response times for most of the test, but there was a warning due to a rare outlier with a maximum response time significantly higher than the 95th percentile. Infrastructure monitoring flagged high memory usage on the load-generator host, which could impact observed metrics if not addressed. No security assessment was provided for this run. Please review outlier latency events and resource consumption during the execution window.

Execution Profile and Artifacts

Execution mode: REAL
Workload profile:
- tool: jmeter
- target_url: https://jsonplaceholder.typicode.com
- endpoint: /posts/1/comments
- method: GET
- users: 15
- max_rps: 10
- duration_seconds: 60
- ramp_up_seconds: 5
- expected_status_code: 200
- p95_threshold_ms: 5000
Artifact directory: src/generated_scripts/Jmeter native reports/JOB-43d9c3eab8984222819f9c2889df9702
Artifact references:
- aligned_time_series: src/generated_scripts/Jmeter native reports/JOB-43d9c3eab8984222819f9c2889df9702/jmeter_PTC_003_plan_aligned_time_series.csv
- generated_script: src/generated_scripts/Jmeter native reports/JOB-43d9c3eab8984222819f9c2889df9702/jmeter_PTC_003_plan.jmx
- jmeter_log: src/generated_scripts/Jmeter native reports/JOB-43d9c3eab8984222819f9c2889df9702/jmeter_PTC_003_plan_jmeter.log
- measurement_metadata: src/generated_scripts/Jmeter native reports/JOB-43d9c3eab8984222819f9c2889df9702/jmeter_PTC_003_plan_measurement_metadata.json

- monitoring_samples: src/generated_scripts/Jmeter native reports/JOB-43d9c3eab8984222819f9c2889df9702/jmeter_PTC_003_plan_runtime_samples.csv
- Native JMeter HTML dashboard
- JMeter dashboard + raw JTL ZIP
- native_report_entrypoint: src/generated_scripts/Jmeter native reports/JOB-43d9c3eab8984222819f9c2889df9702/jmeter_PTC_003_plan_native_report.html
- normalized_metrics: src/generated_scripts/Jmeter native reports/JOB-43d9c3eab8984222819f9c2889df9702/jmeter_PTC_003_plan_normalized_metrics.json
- Raw JMeter results (JTL)
- runner_log: src/generated_scripts/Jmeter native reports/JOB-43d9c3eab8984222819f9c2889df9702/jmeter_PTC_003_plan_runner.log
- time_series: src/generated_scripts/Jmeter native reports/JOB-43d9c3eab8984222819f9c2889df9702/jmeter_PTC_003_plan_time_series.csv

The interactive native HTML dashboard is linked as a separate artifact and is not embedded in this PDF.

Time-Series Charts (REAL Measurements)

Charts show measured values in five-second aligned buckets. Host CPU is from the load-generator machine, not the target.

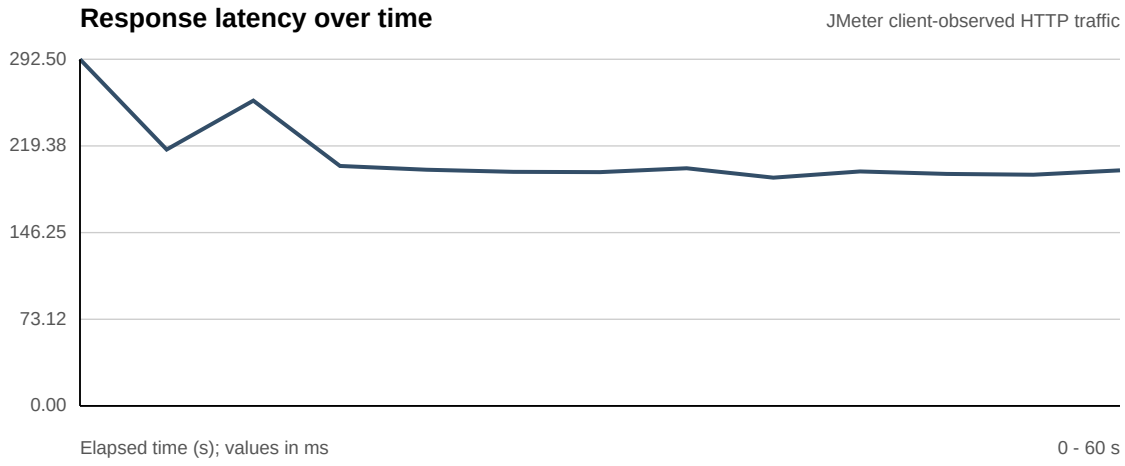

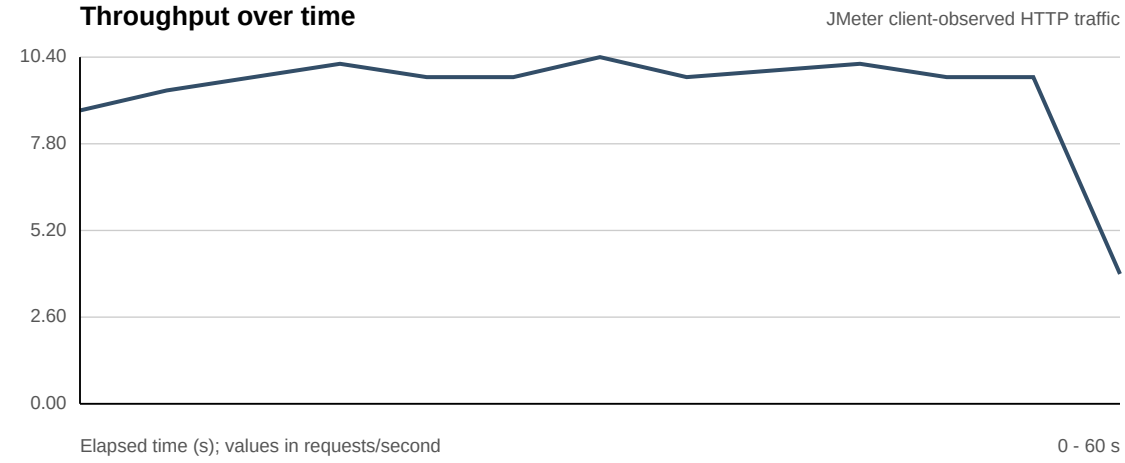

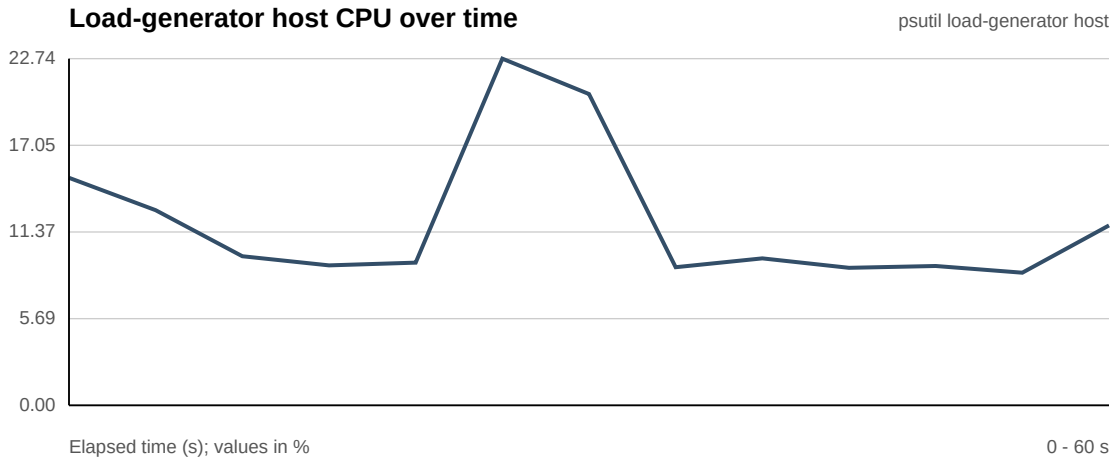

Agent Activity Summary

Agent	*Summary*
Test Execution	Load test executed on /posts/1/comments (Unknown - Content Feedback API — Load Check). Results: 598 total requests, 598 successful, 0 failed. Avg response time: 210.99163879598663 ms; 95th percentile: 257.0 ms; Max: 2970.0 ms. Throughput: 9.597220766939436 TPS; Error rate: 0.0%. Latency percentiles (ms): p50=217.0, p90=246.0, p95=257.0, p99=650.0
Monitoring	CPU usage: 12.0% average, 29.1% peak; memory: 79.43% average, 80.8% peak; DB connections: not collected, TPS: 9.597220766939436, Response time: 210.99163879598663 ms. Data source and scope: REAL: psutil load-generator host samples and jmeter client-observed metrics. Measurement provenance: cpu_usage=psutil (load_generator_host); cpu_peak_usage=psutil (load_generator_host); memory_usage=psutil (load_generator_host); memory_peak_usage=psutil (load_generator_host); response_time=jmeter (load_generator_observed_http_traffic); tps=jmeter (load_generator_observed_http_traffic); db_connections=not_collected (No target-specific database telemetry provider is configured.); total_requests=jmeter JTL (load_generator_observed_http_traffic); failed_requests=jmeter JTL (load_generator_observed_http_traffic); error_rate_pct=jmeter JTL (load_generator_observed_http_traffic); latency_percentiles_ms=jmeter JTL (load_generator_observed_http_traffic). Alerts triggered: ['High load-generator host memory usage']
Security Validation	Security validation was not run for this job.

Agent	*Summary*
Analysis	Test execution completed successfully with no failed requests and healthy response times for most of the test, but there was a warning due to a rare outlier with a maximum response time significantly higher than the 95th percentile. Infrastructure monitoring flagged high memory usage on the load-generator host, which could impact observed metrics if not addressed. No security assessment was provided for this run. Please review outlier latency events and resource consumption during the execution window.
Reporting	A PDF report was generated, including all findings, recommendations, and an executive summary.

Per-API Findings (Test Execution)

- /posts/1/comments [WARNING / medium]
avg=210.99163879598663ms p95=257.0ms max=2970.0ms throughput=9.597220766939436 req/s error_rate=0.0% http_status=200
* Max response time (2970.0ms) is 11.6x p95 (257.0ms) -- outlier spikes present.
history: No historical RAG context available for this API/scenario; this warning is based on threshold breach only.
recommendation: Investigate high max response time outlier. While the median and P95 are healthy, confirm with the team whether rare spikes of this magnitude are expected during load. If not, consult application logs for root cause during the test window.

Infrastructure Health (Monitoring)

Status: WARNING (monitoring_agent reported: WARNING)
CPU usage: 12.0% Memory usage: 79.43% DB connections: None TPS: 9.597220766939436
Response time (monitoring): 210.99163879598663ms
- 1 monitoring alert(s) reported: ['High load-generator host memory usage']
Monitoring alerts: ['High load-generator host memory usage']

Recommended Next Steps

- /posts/1/comments: Investigate high max response time outlier. While the median and P95 are healthy, confirm with the team whether rare spikes of this magnitude are expected during load. If not, consult application logs for root cause during the test window.
- Infrastructure: review CPU/memory/DB connection anomalies above.