Test Case / Original Request

Run PTC-010 with JMeter and gather monitoring results. Ask Security to assess the configured target, then analyze the findings and generate the final PDF with links to the native report and raw data.

REAL Execution Outcome

Execution status: COMPLETED
Runner status: completed
Outcome details: REAL JMeter execution completed. Configured thresholds passed.
Process exit code: 0
Thresholds passed: True

Executive Summary

Application: Unknown
Scenario: Customer Comment API — Load Check
Job ID: JOB-76823ff018244846b6bf67af95080f5c
Overall Result: WARNING

Performance was generally healthy, with low error rate and sustained throughput, but response latency exhibited significant tail spikes—maximum latency was over 4.5 times the p95, consistent with prior results for this endpoint. Infrastructure monitoring reported high memory usage above actionable thresholds on the load-generator host. Security assessment was inconclusive: no live testing was performed due to missing authorized evidence, so results cover only configuration and not in-situ control effectiveness. Address infrastructure resource saturation and monitor tail latency for persistence.

Quick Links

Artifact	*Open*
Native HTML Dashboard	Open →
Dashboard + Raw Results (ZIP)	Open →
Raw Results	Open →

Execution Profile and Artifacts

Execution mode: REAL
Workload profile:
- tool: jmeter
- target_url: https://dummyjson.com
- endpoint: /comments/1
- method: GET
- users: 250
- max_rps: 15

- duration_seconds: 60
- ramp_up_seconds: 5
- expected_status_code: 200
- p95_threshold_ms: 5000
Artifact directory: src/generated_scripts/Jmeter native reports/JOB-76823ff018244846b6bf67af95080f5c
Artifact references:
- aligned_time_series: src/generated_scripts/Jmeter native reports/JOB-76823ff018244846b6bf67af95080f5c/jmeter_PTC_010_plan_aligned_time_series.csv
- generated_script: src/generated_scripts/Jmeter native reports/JOB-76823ff018244846b6bf67af95080f5c/jmeter_PTC_010_plan.jmx
- jmeter_log: src/generated_scripts/Jmeter native reports/JOB-76823ff018244846b6bf67af95080f5c/jmeter_PTC_010_plan_jmeter.log
- measurement_metadata: src/generated_scripts/Jmeter native reports/JOB-76823ff018244846b6bf67af95080f5c/jmeter_PTC_010_plan_measurement_metadata.json
- monitoring_samples: src/generated_scripts/Jmeter native reports/JOB-76823ff018244846b6bf67af95080f5c/jmeter_PTC_010_plan_runtime_samples.csv
- Native JMeter HTML dashboard
- JMeter dashboard + raw JTL ZIP
- native_report_entrypoint: src/generated_scripts/Jmeter native reports/JOB-76823ff018244846b6bf67af95080f5c/jmeter_PTC_010_plan_native_report.html
- normalized_metrics: src/generated_scripts/Jmeter native reports/JOB-76823ff018244846b6bf67af95080f5c/jmeter_PTC_010_plan_normalized_metrics.json
- Raw JMeter results (JTL)
- runner_log: src/generated_scripts/Jmeter native reports/JOB-76823ff018244846b6bf67af95080f5c/jmeter_PTC_010_plan_runner.log
- time_series: src/generated_scripts/Jmeter native reports/JOB-76823ff018244846b6bf67af95080f5c/jmeter_PTC_010_plan_time_series.csv
The interactive native HTML dashboard is linked as a separate artifact and is not embedded in this PDF.

Time-Series Charts (REAL Measurements)

Charts show measured values in five-second aligned buckets. Host CPU is from the load-generator machine, not the target.

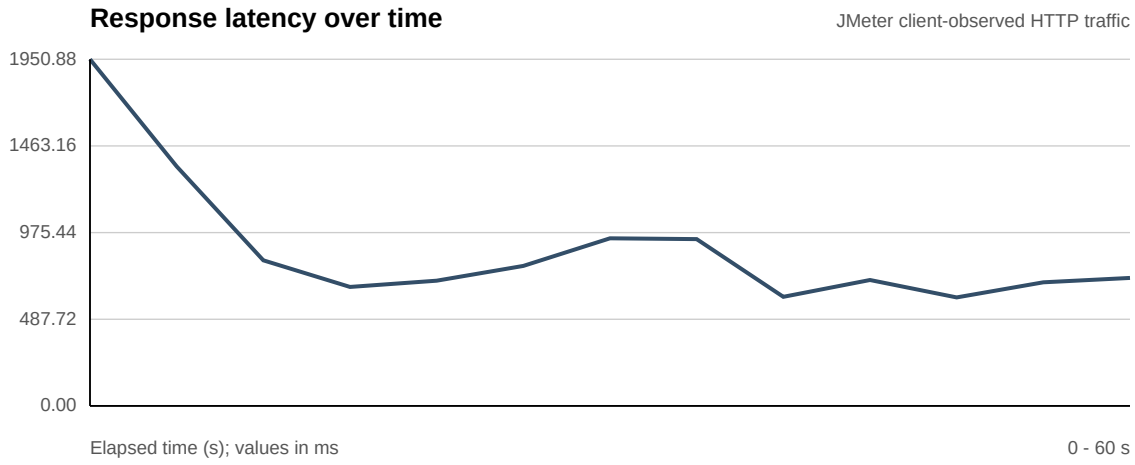

Throughput over time

JMeter client-observed HTTP traffic

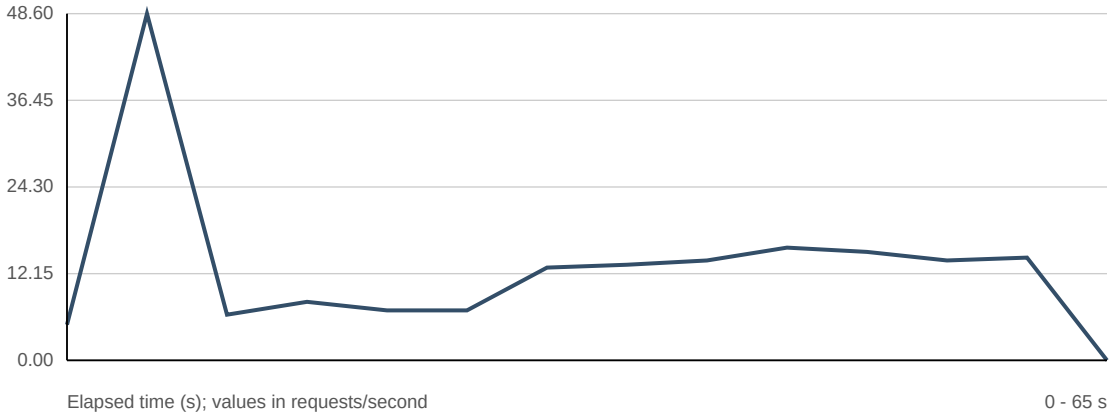

Elapsed time (s); values in requests/second

0 - 65 s

Load-generator host CPU over time

psutil load-generator host

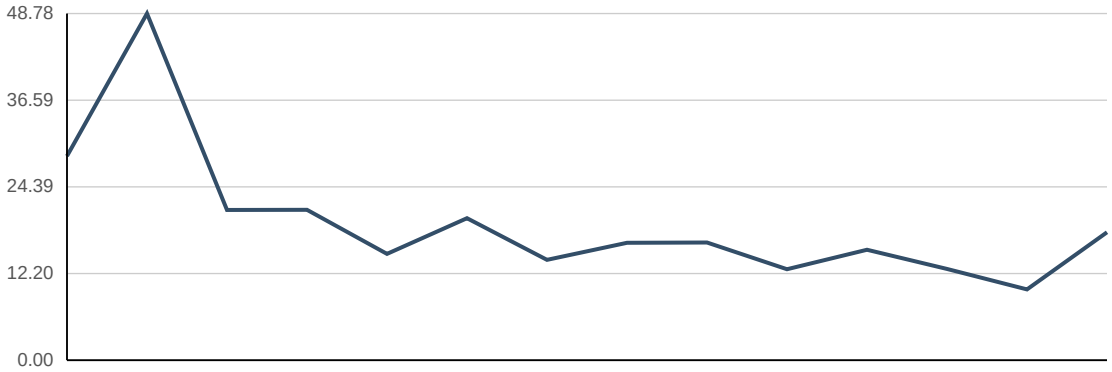

Elapsed time (s); values in %

0 - 65 s

Agent Activity Summary

Agent	*Summary*
Test Execution	Load test executed on /comments/1 (Unknown - Customer Comment API — Load Check). Results: 910 total requests, 910 successful, 0 failed. Avg response time: 936.7153846153847 ms; 95th percentile: 2016.0 ms; Max: 9166.0 ms. Throughput: 13.620885330365 TPS; Error rate: 0.0%. Latency percentiles (ms): p50=628.0, p90=1633.0, p95=2016.0, p99=3581.0

Agent	*Summary*
Monitoring	CPU usage: 19.28% average, 63.7% peak; memory: 87.28% average, 88.3% peak; DB connections: not collected, TPS: 13.620885330365, Response time: 936.7153846153847 ms. Data source and scope: REAL: psutil load-generator host samples and jmeter client-observed metrics. Measurement provenance: cpu_usage=psutil (load_generator_host); cpu_peak_usage=psutil (load_generator_host); memory_usage=psutil (load_generator_host); memory_peak_usage=psutil (load_generator_host); response_time=jmeter (load_generator_observed_http_traffic); tps=jmeter (load_generator_observed_http_traffic); db_connections=not_collected (No target-specific database telemetry provider is configured.); total_requests=jmeter JTL (load_generator_observed_http_traffic); failed_requests=jmeter JTL (load_generator_observed_http_traffic); error_rate_pct=jmeter JTL (load_generator_observed_http_traffic); latency_percentiles_ms=jmeter JTL (load_generator_observed_http_traffic). Alerts triggered: ['High load-generator host memory usage']
Security Validation	Status: WARNING. Security score: N/A. Live security check: not assessed. Issue: The external API surface observation did not complete, so no response-security findings are claimed. Issue: external_surface was not assessed; missing evidence: No authorized external surface observation was completed.. Issue: Scope limitation: configuration validator outcomes describe profile declarations; any live result is a separate single-request observation and does not prove control effectiveness. Issue: Live security check was requested, but no live result was recorded. Control coverage: {'requested_count': 1, 'assessed_count': 0, 'not_assessed_count': 1} Control assessment evidence: - external_surface: not_assessed; basis=live_observation; evidence=[] Assessment limitations: - No authorized external surface observation was completed. - No runtime security checks were performed; results describe configuration declarations only. - The requested live security observation did not complete.
Analysis	Performance was generally healthy, with low error rate and sustained throughput, but response latency exhibited significant tail spikes—maximum latency was over 4.5 times the p95, consistent with prior results for this endpoint. Infrastructure monitoring reported high memory usage above actionable thresholds on the load-generator host. Security assessment was inconclusive: no live testing was performed due to missing authorized evidence, so results cover only configuration and not in-situ control effectiveness. Address infrastructure resource saturation and monitor tail latency for persistence.
Reporting	A PDF report was generated, including all findings, recommendations, and an executive summary.

Per-API Findings (Test Execution)

- /comments/1 [WARNING / medium]

avg=936.7153846153847ms p95=2016.0ms max=9166.0ms throughput=13.620885330365 req/s error_rate=0.0% http_status=200
* Max response time (9166.0ms) is 4.5x p95 (2016.0ms) -- outlier spikes present.
history: This result is consistent with previous runs of the same endpoint, which also showed max and p95 response times significantly above average—prior run (avg: 1317ms, p95: 5158ms, max: 9086ms) exhibited substantial tail latency.
recommendation: Monitor for continued high tail latency, as prior runs for this API also showed these spikes. Investigate database or downstream bottlenecks and review JVM garbage collection if running on Java (see 'Runbook: High Latency on Payments API'). If persistent, consider adding circuit breakers and tightening timeouts. No explicit SLA for this endpoint is documented, but the Payments API SLA and incident postmortem may provide escalation practices.

Infrastructure Health (Monitoring)

Status: WARNING (monitoring_agent reported: WARNING)
CPU usage: 19.28% Memory usage: 87.28% DB connections: None TPS: 13.620885330365 Response time (monitoring): 936.7153846153847ms
- Memory usage 87.28% exceeds 85.0% threshold.
- 1 monitoring alert(s) reported: ['High load-generator host memory usage']
Monitoring alerts: ['High load-generator host memory usage']

Security Findings (Security Agent)

Status: WARNING Security score: N/A
Security result job ID: JOB-76823ff018244846b6bf67af95080f5c.
Live security check: not assessed.
Issues:
- The external API surface observation did not complete, so no response-security findings are claimed.
- external_surface was not assessed; missing evidence: No authorized external surface observation was completed..
- Scope limitation: configuration validator outcomes describe profile declarations; any live result is a separate single-request observation and does not prove control effectiveness.
- Live security check was requested, but no live result was recorded.
Control coverage: {'requested_count': 1, 'assessed_count': 0, 'not_assessed_count': 1}
Requested controls: external_surface
Not assessed controls: external_surface
Control evidence and provenance:
- external_surface: not_assessed; basis=live_observation; evidence fields=[]; references=[]
Assessment limitations:
- No authorized external surface observation was completed.
- No runtime security checks were performed; results describe configuration declarations only.
- The requested live security observation did not complete.

Recommended Next Steps

- /comments/1: Monitor for continued high tail latency, as prior runs for this API also showed these spikes. Investigate database or downstream bottlenecks and review JVM garbage collection if running on Java (see 'Runbook: High Latency on Payments API'). If persistent, consider adding circuit breakers and tightening timeouts. No explicit SLA for this endpoint is documented, but the Payments API SLA and incident postmortem may provide escalation practices.
- Infrastructure: review CPU/memory/DB connection anomalies above.
- Security: address the issues/recommendations listed above before next deployment.