Test Case / Original Request

Run test case PTC-008 using Jmeter, monitor it, analyze the results, and generate the final PDF report with links to the native JMeter report and raw results.

REAL Execution Outcome

Execution status: COMPLETED
Runner status: completed
Outcome details: REAL JMeter execution completed. Configured thresholds passed.
Process exit code: 0
Thresholds passed: True

Executive Summary

Application: Unknown
Scenario: Recipe Details API — Load Check
Job ID: JOB-b2991f8a155444fca3f6b60ed733e82a
Overall Result: WARNING

The Recipe Details API exhibited overall successful execution with no failed requests and stable average response time. However, a warning is noted due to significant max response time outliers, suggesting rare spikes far above the 95th percentile, and the load generator's host showed high memory usage per monitoring alerts. No security evaluation was performed for this run. Investigate request latency spikes and review infrastructure constraints on the load generator.

Execution Profile and Artifacts

Execution mode: REAL
Workload profile:
- tool: jmeter
- target_url: https://dummyjson.com
- endpoint: /recipes/1
- method: GET
- users: 80
- max_rps: 10
- duration_seconds: 60
- ramp_up_seconds: 5
- expected_status_code: 200
- p95_threshold_ms: 5000
Artifact directory: src/generated_scripts/Jmeter native reports/JOB-b2991f8a155444fca3f6b60ed733e82a
Artifact references:
- aligned_time_series: src/generated_scripts/Jmeter native reports/JOB-b2991f8a155444fca3f6b60ed733e82a/jmeter_PTC_008_plan_aligned_time_series.csv
- generated_script: src/generated_scripts/Jmeter native reports/JOB-b2991f8a155444fca3f6b60ed733e82a/jmeter_PTC_008_plan.jmx
- jmeter_log: src/generated_scripts/Jmeter native reports/JOB-b2991f8a155444fca3f6b60ed733e82a/jmeter_PTC_008_plan_jmeter.log
- measurement_metadata: src/generated_scripts/Jmeter native reports/JOB-b2991f8a155444fca3f6b60ed733e82a/jmeter_PTC_008_plan_measurement_metadata.json

- monitoring_samples: src/generated_scripts/Jmeter native reports/JOB-b2991f8a155444fca3f6b60ed733e82a/jmeter_PTC_008_plan_runtime_samples.csv
- Native JMeter HTML dashboard
- JMeter dashboard + raw JTL ZIP
- native_report_entrypoint: src/generated_scripts/Jmeter native reports/JOB-b2991f8a155444fca3f6b60ed733e82a/jmeter_PTC_008_plan_native_report.html
- normalized_metrics: src/generated_scripts/Jmeter native reports/JOB-b2991f8a155444fca3f6b60ed733e82a/jmeter_PTC_008_plan_normalized_metrics.json
- Raw JMeter results (JTL)
- runner_log: src/generated_scripts/Jmeter native reports/JOB-b2991f8a155444fca3f6b60ed733e82a/jmeter_PTC_008_plan_runner.log
- time_series: src/generated_scripts/Jmeter native reports/JOB-b2991f8a155444fca3f6b60ed733e82a/jmeter_PTC_008_plan_time_series.csv

The interactive native HTML dashboard is linked as a separate artifact and is not embedded in this PDF.

Time-Series Charts (REAL Measurements)

Charts show measured values in five-second aligned buckets. Host CPU is from the load-generator machine, not the target.

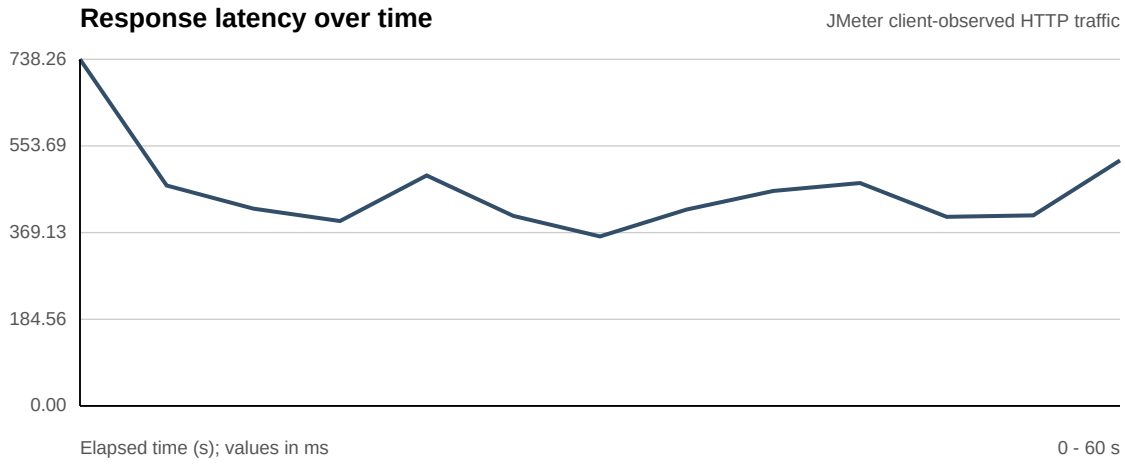

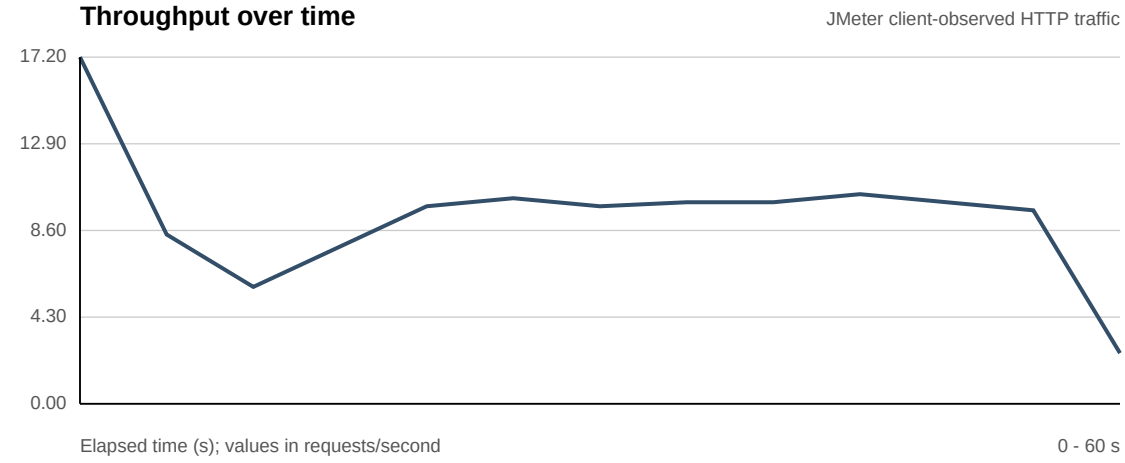

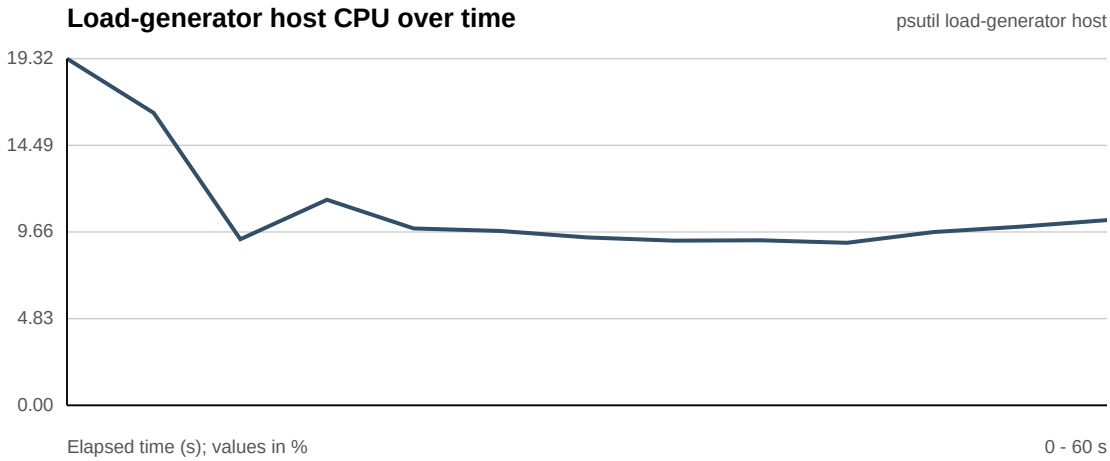

Agent Activity Summary

Agent	*Summary*
Test Execution	Load test executed on /recipes/1 (Unknown - Recipe Details API — Load Check). Results: 604 total requests, 604 successful, 0 failed. Avg response time: 473.3841059602649 ms; 95th percentile: 974.0 ms; Max: 3666.0 ms. Throughput: 9.500275880956154 TPS; Error rate: 0.0%. Latency percentiles (ms): p50=362.0, p90=626.0, p95=974.0, p99=2301.0
Monitoring	CPU usage: 10.85% average, 34.0% peak; memory: 79.8% average, 81.8% peak; DB connections: not collected, TPS: 9.500275880956154, Response time: 473.3841059602649 ms. Data source and scope: REAL: psutil load-generator host samples and jmeter client-observed metrics. Measurement provenance: cpu_usage=psutil (load_generator_host); cpu_peak_usage=psutil (load_generator_host); memory_usage=psutil (load_generator_host); memory_peak_usage=psutil (load_generator_host); response_time=jmeter (load_generator_observed_http_traffic); tps=jmeter (load_generator_observed_http_traffic); db_connections=not_collected (No target-specific database telemetry provider is configured.); total_requests=jmeter JTL (load_generator_observed_http_traffic); failed_requests=jmeter JTL (load_generator_observed_http_traffic); error_rate_pct=jmeter JTL (load_generator_observed_http_traffic); latency_percentiles_ms=jmeter JTL (load_generator_observed_http_traffic). Alerts triggered: ['High load-generator host memory usage']
Security Validation	Security validation was not run for this job.
Analysis	The Recipe Details API exhibited overall successful execution with no failed requests and stable average response time. However, a warning is noted due to significant max response time outliers, suggesting rare spikes far above the 95th percentile, and the load generator's host showed high memory usage per monitoring alerts. No security evaluation was performed for this run. Investigate request latency spikes and review infrastructure constraints on the load generator.

Agent	*Summary*
Reporting	A PDF report was generated, including all findings, recommendations, and an executive summary.

Per-API Findings (Test Execution)

- /recipes/1 [WARNING / medium]
avg=473.3841059602649ms p95=974.0ms max=3666.0ms throughput=9.500275880956154 req/s error_rate=0.0% http_status=200
* Max response time (3666.0ms) is 3.8x p95 (974.0ms) -- outlier spikes present.
history: No historical context was available; relying on static threshold breach for analysis.
recommendation: Investigate the response time spikes (max response time is significantly above the 95th percentile). Since no historical data, focus remediation on request tracing and high-latency outlier handling.

Infrastructure Health (Monitoring)

Status: WARNING (monitoring_agent reported: WARNING)
CPU usage: 10.85% Memory usage: 79.8% DB connections: None TPS: 9.500275880956154
Response time (monitoring): 473.3841059602649ms
- 1 monitoring alert(s) reported: ['High load-generator host memory usage']
Monitoring alerts: ['High load-generator host memory usage']

Recommended Next Steps

- /recipes/1: Investigate the response time spikes (max response time is significantly above the 95th percentile). Since no historical data, focus remediation on request tracing and high-latency outlier handling.
- Infrastructure: review CPU/memory/DB connection anomalies above.