Test Case / Original Request

Start a real k6 run for PTC-003. Monitor it, assess security for the same target, analyze the current run’s results, and generate a PDF report with links to its native report and raw data.

REAL Execution Outcome

Execution status: COMPLETED
Runner status: completed
Outcome details: REAL k6 execution completed. Threshold breaches, if any, are retained as performance findings.
Thresholds passed: True

Executive Summary

Application: Unknown
Scenario: Content Feedback API — Load Check
Job ID: JOB-8472a391bda441c197e2aee617876acf
Overall Result: WARNING

The load test completed successfully, with no failed requests and steady throughput. However, significant latency variance was observed (p95 latency at 3.6x the average, and max latency 2.2x the p95), suggesting some tail latency spikes, although all requests succeeded. Infrastructure resource usage was normal and there were no infra anomalies. Security assessment flagged a warning: the HSTS header was not observed for this API, which should be reviewed for HTTPS best practices. No runbooks or SLA matches were available for these findings.

Execution Profile and Artifacts

Execution mode: REAL
Workload profile:
- tool: k6
- target_url: https://jsonplaceholder.typicode.com
- endpoint: /posts/1/comments
- method: GET
- users: 15
- max_rps: 10
- duration_seconds: 60
- ramp_up_seconds: 5
- expected_status_code: 200
- p95_threshold_ms: 5000
Artifact directory: src/generated_scripts/k6 native reports/JOB-8472a391bda441c197e2aee617876acf
Artifact references:
- aligned_time_series: src/generated_scripts/k6 native reports/JOB-8472a391bda441c197e2aee617876acf/k6_PTC_003_Content_Feedback_API_Load_Check_aligned_time_series.csv
- generated_script: src/generated_scripts/k6 native reports/JOB-8472a391bda441c197e2aee617876acf/k6_PTC_003_Content_Feedback_API_Load_Check.js
- k6_summary: src/generated_scripts/k6 native reports/JOB-8472a391bda441c197e2aee617876acf/k6_PTC_003_Content_Feedback_API_Load_Check_summary.json
- measurement_metadata: src/generated_scripts/k6 native reports/JOB-8472a391bda441c197e2aee617876acf/k6_PTC_003_Content_Feedback_API_Load_Check_measurement_metadata.json

- monitoring_samples: src/generated_scripts/k6 native reports/JOB-8472a391bda441c197e2aee617876acf/k6_PTC_003_Content_Feedback_API_Load_Check_runtime_samples.csv
- Native k6 HTML report
- Native report + raw results ZIP
- normalized_metrics: src/generated_scripts/k6 native reports/JOB-8472a391bda441c197e2aee617876acf/k6_PTC_003_Content_Feedback_API_Load_Check_normalized_metrics.json
- Raw k6 results (JSON)
- runner_log: src/generated_scripts/k6 native reports/JOB-8472a391bda441c197e2aee617876acf/k6_PTC_003_Content_Feedback_API_Load_Check_runner.log
- time_series: src/generated_scripts/k6 native reports/JOB-8472a391bda441c197e2aee617876acf/k6_PTC_003_Content_Feedback_API_Load_Check_time_series.csv

The interactive native HTML dashboard is linked as a separate artifact and is not embedded in this PDF.

Time-Series Charts (REAL Measurements)

Charts show measured values in five-second aligned buckets. Host CPU is from the load-generator machine, not the target.

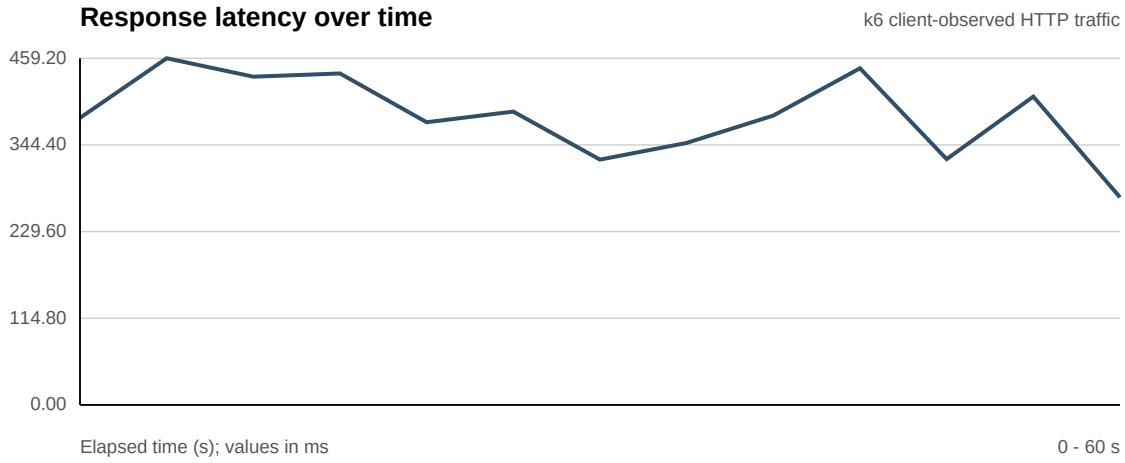

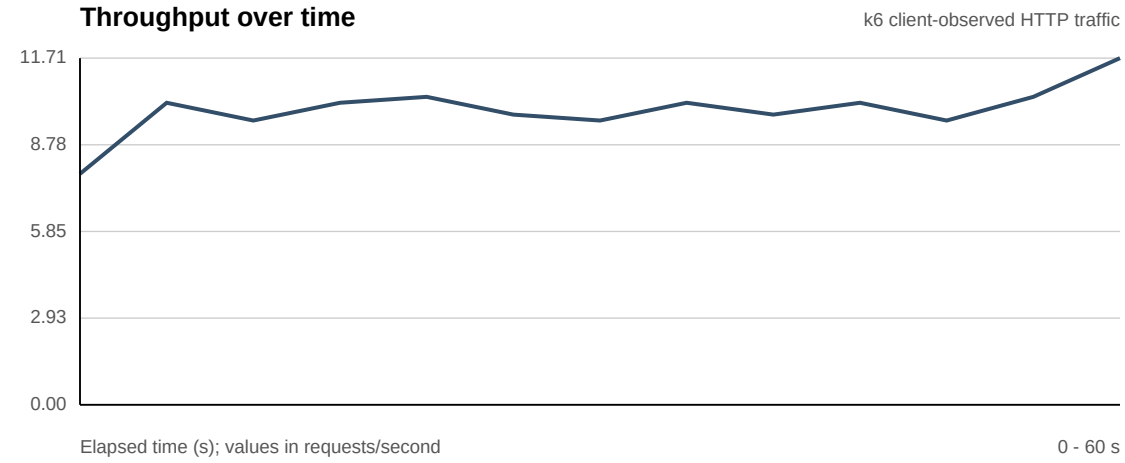

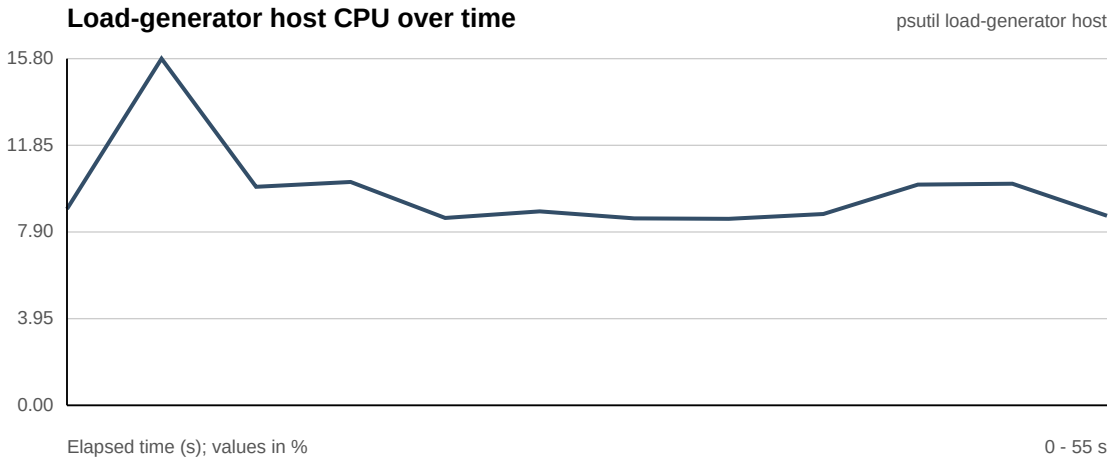

Agent Activity Summary

Agent	*Summary*
Test Execution	Load test executed on /posts/1/comments (Unknown - Content Feedback API — Load Check). Results: 590 total requests, 590 successful, 0 failed. Avg response time: 393.188994915254 ms; 95th percentile: 1399.4311499999988 ms; Max: 3019.649 ms. Throughput: 9.827247810126913 TPS; Error rate: 0.0%.
Monitoring	CPU usage: 9.75% average, 25.9% peak; memory: 67.86% average, 69.7% peak; DB connections: not collected, TPS: 9.827247810126913, Response time: 393.188994915254 ms. Data source and scope: REAL: psutil load-generator host samples and k6 client-observed metrics. Measurement provenance: cpu_usage=psutil (load_generator_host); cpu_peak_usage=psutil (load_generator_host); memory_usage=psutil (load_generator_host); memory_peak_usage=psutil (load_generator_host); response_time=k6 (load_generator_observed_http_traffic); tps=k6 (load_generator_observed_http_traffic); db_connections=not_collected (No target-specific database telemetry provider is configured.). No alerts triggered.

Agent	*Summary*
Security Validation	Status: WARNING. Security score: 0%. Live security check: completed. Issue: Strict-Transport-Security was not present in the observed response; this single response does not establish the full origin policy. Issue: Scope limitation: configuration validator outcomes describe profile declarations; any live result is a separate single-request observation and does not prove control effectiveness. Issue: Live HEAD observation completed with HTTP status 200; response headers are observations only, not proof of control effectiveness. Control coverage: {'requested_count': 1, 'assessed_count': 1, 'not_assessed_count': 0} Control assessment evidence: - external_surface: assessed; basis=live_observation; evidence=['live-observation:JOB-8472a391bda441c197e2aee617876acf'] Evidence references: live-observation:JOB-8472a391bda441c197e2aee617876acf Recommendation: Review whether an HSTS policy is appropriate for this HTTPS API origin.
Analysis	The load test completed successfully, with no failed requests and steady throughput. However, significant latency variance was observed (p95 latency at 3.6x the average, and max latency 2.2x the p95), suggesting some tail latency spikes, although all requests succeeded. Infrastructure resource usage was normal and there were no infra anomalies. Security assessment flagged a warning: the HSTS header was not observed for this API, which should be reviewed for HTTPS best practices. No runbooks or SLA matches were available for these findings.
Reporting	A PDF report was generated, including all findings, recommendations, and an executive summary.

Per-API Findings (Test Execution)

- /posts/1/comments [WARNING / high]
avg=393.188994915254ms p95=1399.4311499999988ms max=3019.649ms throughput=9.827247810126913 req/s error_rate=0.0% http_status=200
* p95 (1399.4311499999988ms) is 3.6x the average (393.188994915254ms) -- wide latency variance / possible tail contention.
* Max response time (3019.649ms) is 2.2x p95 (1399.4311499999988ms) -- outlier spikes present.
history: No RAG-based historical context available; threshold breach is interpreted based only on this run's measured values.
recommendation: Review if high p95/max latency is expected and investigate cause of tail latency or outliers if not. No runbook, postmortem, or explicit SLA match fetched for this API, as historical context was unavailable.

Infrastructure Health (Monitoring)

Status: PASS (monitoring_agent reported: OK)
CPU usage: 9.75% Memory usage: 67.86% DB connections: None TPS: 9.827247810126913
Response time (monitoring): 393.188994915254ms

Security Findings (Security Agent)

Status: WARNING Security score: 0%
Security result job ID: JOB-8472a391bda441c197e2aee617876acf.
Live security check: completed.
Issues:
- Strict-Transport-Security was not present in the observed response; this single response does not establish the full origin policy.
- Scope limitation: configuration validator outcomes describe profile declarations; any live result is a separate single-request observation and does not prove control effectiveness.
- Live HEAD observation completed with HTTP status 200; response headers are observations only, not proof of control effectiveness.
Control coverage: {'requested_count': 1, 'assessed_count': 1, 'not_assessed_count': 0}
Requested controls: external_surface
Assessed controls: external_surface
Control evidence and provenance:
- external_surface: assessed; basis=live_observation; evidence fields=['target_origin', 'endpoint', 'method', 'response_status_code', 'observed_security_headers']; references=['live-observation:JOB-8472a391bda441c197e2aee617876acf']
Evidence references: live-observation:JOB-8472a391bda441c197e2aee617876acf
Live observation: HEAD /posts/1/comments; HTTP 200; elapsed 324.55245906021446 ms; requests sent 1.
Observed response headers (observations only):
- cache-control: max-age=43200
- x-content-type-options: nosniff
Recommendations:
- Review whether an HSTS policy is appropriate for this HTTPS API origin.

Recommended Next Steps

- /posts/1/comments: Review if high p95/max latency is expected and investigate cause of tail latency or outliers if not. No runbook, postmortem, or explicit SLA match fetched for this API, as historical context was unavailable.
- Security: address the issues/recommendations listed above before next deployment.